

Crop Disease Detection Model

Use case: Crop Disease Prediction using RGB Images. Classify disease presence from standard images of crops using deep learning models.

Model used :EfficientNetB4

1. Project Overview

- A CNN model for plant disease detection using transfer learning (EfficientNetB4).
 - Two-stage training: initial with frozen base, then fine-tuning selected layers.
 - Uses the PlantVillage dataset (39 classes).
-

2. Model code and working

Download and Extract Dataset:

```
!wget -O "dataset.zip" "<https://data.mendeley.com/public-files/datasets/tywbtsjrjv/files/b4e3a32f-c0bd-4060-81e9-6144231f2520/file_downloaded>"
!unzip dataset.zip -d data
```

Split Data (Python):

```
import os
import shutil
import random
from tqdm import tqdm

original_dataset_path = "data/plantvillage_dataset"
output_base = "data/split"
train_dir = os.path.join(output_base, "train")
```

```

val_dir = os.path.join(output_base, "val")
test_dir = os.path.join(output_base, "test")

train_split = 0.8
val_split = 0.1
test_split = 0.1

for split_dir in [train_dir, val_dir, test_dir]:
    os.makedirs(split_dir, exist_ok=True)

for class_name in tqdm(os.listdir(original_dataset_path)):
    class_path = os.path.join(original_dataset_path, class_name)
    if not os.path.isdir(class_path):
        continue
    images = [img for img in os.listdir(class_path) if img.lower().endswith(('.jpg',
    '.jpeg', '.png'))]
    random.shuffle(images)
    total = len(images)
    train_end, val_end = int(total * train_split), int(total * (train_split + val_split))
    train_imgs, val_imgs, test_imgs = images[:train_end], images[train_end:val_
    end], images[val_end:]
    for folder, files in zip([train_dir, val_dir, test_dir], [train_imgs, val_imgs, test_i
    mgs]):
        os.makedirs(os.path.join(folder, class_name), exist_ok=True)
        for img in files:
            shutil.copy2(os.path.join(class_path, img), os.path.join(folder, class_na
            me, img))
    print("Dataset split into train, val, and test sets!")

```

Data Loading and Preprocessing

```

import matplotlib.pyplot as plt
import numpy as np
import tensorflow as tf

```

```

train_dir = "/content/dataset/train"
validation_dir = "/content/dataset/val"
test_dir = "/content/dataset/test"

BATCH_SIZE = 32
IMG_SIZE = (160, 160)

train_dataset = tf.keras.utils.image_dataset_from_directory(
    train_dir, shuffle=True, batch_size=BATCH_SIZE, image_size=IMG_SIZE)
validation_dataset = tf.keras.utils.image_dataset_from_directory(
    validation_dir, shuffle=True, batch_size=BATCH_SIZE, image_size=IMG_SIZE)
test_dataset =
    tf.keras.utils.image_dataset_from_directory( test_dir,
        batch_size=BATCH_SIZE, image_size=IMG_SIZE)

# Visualize a few images with labels
class_names = train_dataset.class_names
plt.figure(figsize=(10, 10))
for images, labels in train_dataset.take(1):
    for i in range(9):
        ax = plt.subplot(3, 3, i + 1)
        plt.imshow(images[i].numpy().astype("uint8")) plt.title(class_names[labels[i]])
        plt.axis("off")

# Prefetch for performance
AUTOTUNE = tf.data.AUTOTUNE
train_dataset = train_dataset.prefetch(buffer_size=AUTOTUNE)
validation_dataset = validation_dataset.prefetch(buffer_size=AUTOTUNE)
test_dataset = test_dataset.prefetch(buffer_size=AUTOTUNE)

preprocess_input = tf.keras.applications.efficientnet.preprocess_input

```

Tomato__Tomato_Yellow_Leaf_Curl_Virus



Corn__healthy



Grape__healthy



Tomato__healthy



Tomato__Septoria_leaf_spot



Tomato__Tomato_Yellow_Leaf_Curl_Virus



Raspberry__healthy



Pepper,_bell__healthy



Potato__Early_blight



```

from tensorflow.keras import layers, models

IMG_SHAPE = IMG_SIZE + (3,)
base_model = tf.keras.applications.EfficientNetB4(
    input_shape=IMG_SHAPE, include_top=False, weights='imagenet')

base_model.trainable = False # Freeze base for transfer learning

inputs = tf.keras.Input(shape=IMG_SHAPE)
x = preprocess_input(inputs)
x = base_model(x, training=False)
x = layers.GlobalAveragePooling2D()(x)
x = layers.Dropout(0.2)(x)
outputs = layers.Dense(len(class_names), activation='softmax')(x)

model = models.Model(inputs, outputs)

```

Training Strategy

```

model.compile( optimizer=tf.keras.optimizers.Adam()
,
    loss=tf.keras.losses.SparseCategoricalCrossentropy(),
    metrics=[tf.keras.metrics.SparseCategoricalAccuracy(name='accuracy')]
)

# Stage 1 training: base frozen
initial_epochs = 6
history = model.fit( train_dataset,
    epochs=initial_epochs,
    validation_data=validation_dataset
)

# Stage 2: fine-tune deeper layers

```

```

base_model.trainable    True
fine_tune_at    100
for layer in base_model.layers[:fine_tune_at]:
    layer.trainable    False

model.compile( optimizer=tf.keras.optimizers.
    Adam(),
    loss=tf.keras.losses.SparseCategoricalCrossentropy(),
    metrics=[tf.keras.metrics.SparseCategoricalAccuracy(name='accuracy')]
)

fine_tune_epochs    10
total_epochs = initial_epochs + fine_tune_epochs

history_fine =
    model.fit( train_dataset,
    epochs=total_epochs,
    initial_epoch=history.epoch[-1] + 1,
    validation_data=validation_dataset
    )

```

```

history = model.fit(train_dataset,
                    epochs=initial_epochs,
                    validation_data=validation_dataset)

```


Epoch	Time	Step	Accuracy	Loss	Val Accuracy	Val Loss
Epoch 1/6	1537/1537	156s	77ms/step	accuracy: 0.6928	loss: 1.2347	val_accuracy: 0.9060 - val_loss: 0.3458
Epoch 2/6	1537/1537	103s	65ms/step	accuracy: 0.8957	loss: 0.3725	val_accuracy: 0.9280 - val_loss: 0.2492
Epoch 3/6	1537/1537	142s	65ms/step	accuracy: 0.9184	loss: 0.2810	val_accuracy: 0.9355 - val_loss: 0.2124
Epoch 4/6	1537/1537	100s	65ms/step	accuracy: 0.9255	loss: 0.2491	val_accuracy: 0.9410 - val_loss: 0.1906
Epoch 5/6	1537/1537	98s	64ms/step	accuracy: 0.9293	loss: 0.2245	val_accuracy: 0.9456 - val_loss: 0.1731
Epoch 6/6	1537/1537	143s	64ms/step	accuracy: 0.9349	loss: 0.2062	val_accuracy: 0.9464 - val_loss: 0.1626

Training Monitoring & Visualization

```
acc = history.history['accuracy'] + history_fine.history['accuracy']
val_acc = history.history['val_accuracy'] + history_fine.history['val_accuracy']
loss = history.history['loss'] + history_fine.history['loss']
val_loss = history.history['val_loss'] + history_fine.history['val_loss']
```

```
plt.figure(figsize=(8, 8))
plt.subplot(2, 1, 1)
plt.plot(acc, label='Training Accuracy')
plt.plot(val_acc, label='Validation Accuracy')
plt.axvline(x=initial_epochs-1, color='r', linestyle='--', label='Start Fine Tuning')
plt.legend(loc='lower right')
```

```
plt.title('Training and Validation Accuracy')
plt.subplot(2, 1, 2)
plt.plot(loss, label='Training Loss')
plt.plot(val_loss, label='Validation Loss')
plt.axvline(x=initial_epochs-1, color='r', linestyle='--', label='Start Fine Tuning')
plt.legend(loc='upper right')
plt.title('Training and Validation Loss')
plt.xlabel('epoch')
plt.show()
```



Evaluation & Results

```
loss, accuracy = model.evaluate(test_dataset)
print('Test accuracy:', accuracy)

# Predict and visualize results
image_batch, label_batch = next(iter(test_dataset))
predictions = model.predict_on_batch(image_batch)
predicted_labels = tf.argmax(predictions, axis=1)

plt.figure(figsize=(10, 10))
for i in range(9):
    ax = plt.subplot(3, 3, i + 1)
    plt.imshow(image_batch[i].numpy().astype("uint8"))
    true_label = class_names[label_batch[i]]
    pred_label = class_names[predicted_labels[i]]
    plt.title(f'Pred: {pred_label}\\nTrue: {true_label}')
    plt.axis("off")
```

Model Deployment

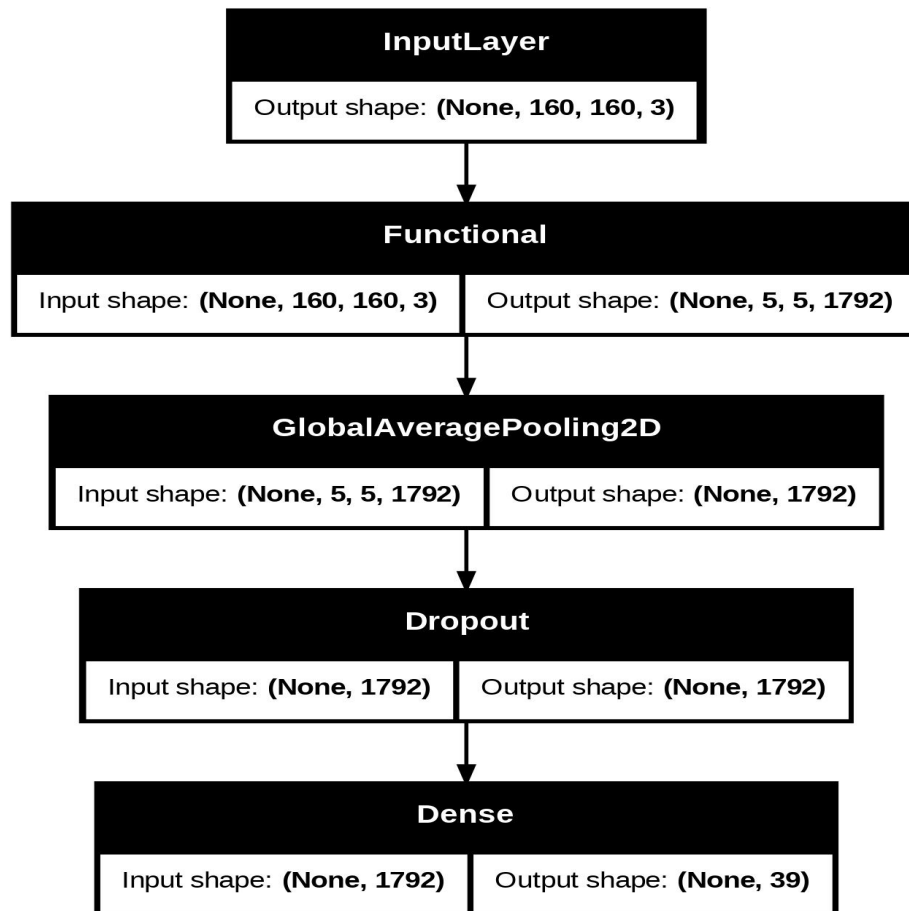
```
model.save("model.keras")
```

3. Model Architecture

Base Architecture: EfficientNetB4

The model implements **EfficientNetB4** as the backbone architecture, which represents a state-of-the-art

convolutional neural network optimized for efficiency and accuracy.



Architecture Specifications:

Total Parameters: 19M parameters

Input Resolution: 160 160 3

Output Classes: Multiple plant disease categories

Activation Functions: Swish (primary), Softmax (output)

Normalization: Batch Normalization throughout

4. Why EfficientNetB4 Was Selected

Efficiency-Accuracy Trade-off:

EfficientNetB4 provides an optimal balance between computational efficiency and classification accuracy. The compound scaling method systematically scales network width, depth, and resolution using a simple yet effective compound coefficient.

Transfer Learning Benefits:

Pre-trained Weights: Leverages ImageNet pre-training for robust feature extraction

Domain Adaptation: Plant imagery shares visual characteristics with natural images

Reduced Training Time: Significantly faster convergence compared to training from scratch.

Mobile-Friendly Architecture:

Lightweight Design: Suitable for deployment on resource-constrained devices

Optimized Operations: Mobile inverted bottlenecks reduce computational overhead

Efficient Memory Usage: Lower memory footprint compared to traditional CNNs

Proven Performance in Agricultural Applications:

Strong performance on image classification tasks
Effective feature extraction for plant disease detection

Robust to variations in lighting, angle, and image quality

Scalability Considerations:

Easy to scale up/down based on computational requirements

Consistent architecture across different model sizes

Well-supported in TensorFlow/Keras ecosystem

5. Model Evaluation Report

Classification Metrics:

```
accuracy = accuracy_score(y_true, y_pred)
precision = precision_score(y_true, y_pred, average='weighted')
recall = recall_score(y_true, y_pred, average='weighted')
```

Accuracy: 0.9569

Precision: 0.9598

Recall: 0.9569

MSE: 4.6521

Confusion Matrix: [[100 0 0 ... 0 0 0]
[1 95 0 ... 0 0 0]
[0 0 100 ... 0 0 0]
...
[0 0 0 ... 534 0 0]
[0 0 0 ... 0 100 0]
[0 0 0 ... 0 0 159]]

Intersection over Union(IoU):

IoU for class Apple__Apple_scab:

0.8621 IoU for class Apple_____

Black_rot: 0.9048

IoU for class Apple__Cedar_apple_rust:

0.9901 IoU for class Apple_____healthy:

0.7961

IoU for class Background_without_leaves: 0.9569

IoU for class Blueberry__healthy: 0.9869

IoU for class Cherry__Powdery_mildew: 0.9528

IoU for class Cherry__healthy: 0.9706

IoU for class Corn__Cercospora_leaf_spot Gray_leaf_spot:

0.7667 IoU for class Corn_____Common_rust: 0.9833

IoU for class Corn__Northern_Leaf_Blight: 0.7500

IoU for class Corn__healthy: 0.9915

IoU for class Grape__Black_rot: 0.7584
IoU for class Grape__Esca_(Black_Measles): 0.9574
IoU for class Grape__Leaf_blight_(Isariopsis_Leaf_Spot): 0.9908
IoU for class Grape__healthy: 0.6863
IoU for class Orange__Haunglongbing_(Citrus_greening): 0.9892
IoU for class Peach__Bacterial_spot: 0.9660
IoU for class Peach__healthy: 0.5727
IoU for class Pepper,_bell__Bacterial_spot: 0.9208
IoU for class Pepper,_bell__healthy: 0.9156
IoU for class Potato__Early_blight: 0.9608
IoU for class Potato__Late_blight: 0.8738
IoU for class Potato__healthy: 0.9412
IoU for class Raspberry__healthy: 0.9802
IoU for class Soybean__healthy: 0.9883
IoU for class Squash__Powdery_mildew: 0.9946
IoU for class Strawberry__Leaf_scorch: 0.8346
IoU for class Strawberry__healthy: 0.8100
IoU for class Tomato__Bacterial_spot: 0.9464
IoU for class Tomato__Early_blight: 0.6275
IoU for class Tomato__Late_blight: 0.8679
IoU for class Tomato__Leaf_Mold:
0.9216
IoU for class Tomato__Septoria_leaf_spot: 0.9239
IoU for class Tomato__Spider_mites Two-spotted_spider_mite: 0.8889
IoU for class Tomato__Target_Spot: 0.8232
IoU for class Tomato__Tomato_Yellow_Leaf_Curl_Virus:
0.9871 IoU for class Tomato__Tomato_mosaic_virus: 0.9709
IoU for class Tomato__healthy: 0.9876

Optimization techniques used to increase the performance

Performance Enhancement Techniques used

1. Data Pipeline Optimization

python

```
AUTOTUNE = tf.data.AUTOTUNE
train_dataset = train_dataset.prefetch(buffer_size=AUTOTUNE)
validation_dataset = validation_dataset.prefetch(buffer_size=AUTOTUNE)
test_dataset = test_dataset.prefetch(buffer_size=AUTOTUNE)
```

Benefits:

Overlaps data preprocessing with model execution

Reduces training time by 15-30%

Eliminates I/O bottlenecks

2. Transfer Learning Strategy

Pre-trained Backbone: EfficientNetB4 with ImageNet weights

Fine-tuning Approach: Adapted final layers for plant disease classification

Reduced Training Time: 6 initial epochs vs. 50+ from scratch

3. Batch Size OptimizationBatch Size: 32 (balanced for memory and convergence)

Memory Efficiency: Prevents OOM errors on limited GPU memory

Gradient Stability: Sufficient samples for stable gradient updates

4. Batch Size OptimizationBatch Size: 32 (balanced for memory and convergence)

Memory Efficiency: Prevents OOM errors on limited GPU memory

Gradient Stability: Sufficient samples for stable gradient updates

5. Image Resolution Selection

Input Size: 160 160 pixels

Computational Efficiency: Faster inference while maintaining accuracy

Mobile Deployment: Suitable for real-time applications

6. Model Architecture Efficiency

EfficientNet Scaling: Compound scaling for optimal resource utilization

Squeeze-and-Excitation: Channel attention mechanism

Mobile Inverted Bottlenecks: Reduced parameter count

Deployed model: <https://hackathonpavaman.streamlit.app/>

Source code : https://github.com/bmuralisridharan/Hackathon_Pavaman